

EXPERIMENT NO – 6

Build a Convolution Neural Network for simple image (dogs and Cats) Classification

Image classification is the task of categorizing images into different classes based on their content. In this case, we want to build a model that can distinguish between images of dogs and cats.

Program:

```
# Install unrar if not installed
```

```
!apt-get install unrar -y
```

```
# Extract the dataset
```

```
!unrar x "/content/dogs-vs-cats.rar" "/content/dogs-vs-cats/"
```

```
# Verify extraction
```

```
!ls "/content/dogs-vs-cats/"
```

```
import tensorflow as tf
```

```
from tensorflow.keras.preprocessing.image import ImageDataGenerator
```

```
# Define dataset directory
```

```
dataset_dir = "/content/dogs-vs-cats/"
```

```
# Define image size & batch size
```

```
IMG_SIZE = (150, 150)
```

```
BATCH_SIZE = 32
```

```
# Data augmentation for training set
```

```
train_datagen = ImageDataGenerator(
```

```
    rescale=1.0/255,
```

```
    rotation_range=20,
```

```
    width_shift_range=0.2,
```

```
    height_shift_range=0.2,
```

```
    shear_range=0.2,
```

```

        zoom_range=0.2,
        horizontal_flip=True,
        validation_split=0.2 # 80% train, 20% validation
    )

# Load training & validation datasets
train_generator = train_datagen.flow_from_directory(
    dataset_dir,
    target_size=IMG_SIZE,
    batch_size=BATCH_SIZE,
    class_mode='binary', # Since it's a binary classification (dogs vs cats)
    subset='training'
)

val_generator = train_datagen.flow_from_directory(
    dataset_dir,
    target_size=IMG_SIZE,
    batch_size=BATCH_SIZE,
    class_mode='binary',
    subset='validation'
)

Found 4 images belonging to 1 classes.
Found 1 images belonging to 1 classes.

from tensorflow.keras.models import Sequential

from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Dropout

# Build the CNN Model
model = Sequential([
    Conv2D(32, (3,3), activation='relu', input_shape=(150, 150, 3), strides=(1,1), padding="same"),
    MaxPooling2D((2,2), strides=(2,2)),

```

```

Conv2D(64, (3,3), activation='relu', padding="same"),
MaxPooling2D((2,2), strides=(2,2)),

Conv2D(128, (3,3), activation='relu', padding="same"),
MaxPooling2D((2,2), strides=(2,2)),

Flatten(),
Dense(512, activation='relu'),
Dropout(0.5),
Dense(1, activation='sigmoid') # Binary classification (dog or cat)
])

# Compile the model

model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])

```

Model Summary

```
model.summary()
```

```

/usr/local/lib/python3.11/dist-
packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do
not pass an `input_shape`/`input_dim` argument to a layer. When using
Sequential models, prefer using an `Input(shape)` object as the first layer
in the model instead.

```

```

super().__init__(activity_regularizer=activity_regularizer, **kwargs)
Model: "sequential"

```

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 150, 150, 32)	896
max_pooling2d (MaxPooling2D)	(None, 75, 75, 32)	0
conv2d_1 (Conv2D)	(None, 75, 75, 64)	18,496
max_pooling2d_1 (MaxPooling2D)	(None, 37, 37, 64)	0
conv2d_2 (Conv2D)	(None, 37, 37, 128)	73,856
max_pooling2d_2 (MaxPooling2D)	(None, 18, 18, 128)	0
flatten (Flatten)	(None, 41472)	0
dense (Dense)	(None, 512)	21,234,176
dropout (Dropout)	(None, 512)	0
dense_1 (Dense)	(None, 1)	513

```
Total params: 21,327,937 (81.36 MB)
Trainable params: 21,327,937 (81.36 MB)
Non-trainable params: 0 (0.00 B)
# Train the model
```

```
history = model.fit(
    train_generator,
    validation_data=val_generator,
    epochs=10
)
```

```
/usr/local/lib/python3.11/dist-packages/keras/src/trainers/data_adapters/py_dataset_adapter.py:121:
UserWarning: Your `PyDataset` class should call `super().__init__(**kwargs)` in its constructor. `**kwargs`
can include `workers`, `use_multiprocessing`, `max_queue_size`. Do not pass these arguments to `fit()`, as
they will be ignored.
```

```
self._warn_if_super_not_called()
```

```
Epoch 1/10
```

```
1/1 ----- 5s 5s/step - accuracy: 0.2500 - loss: 0.7157 -
```

```
val_accuracy: 1.0000 - val_loss: 2.0281e-10
```

```
Epoch 2/10
```

```
1/1 ----- 1s 1s/step - accuracy: 1.0000 - loss: 1.9600e-06 -
```

```
val_accuracy: 1.0000 - val_loss: 1.3092e-19
```

```
Epoch 3/10
```

```
1/1 ----- 1s 871ms/step - accuracy: 1.0000 - loss: 2.0541e-09 -
```

```
val_accuracy: 1.0000 - val_loss: 3.1983e-30
```

```
Epoch 4/10
```

```
1/1 ----- 1s 815ms/step - accuracy: 1.0000 - loss: 3.4824e-15 -
```

```
val_accuracy: 1.0000 - val_loss: 0.0000e+00
```

```
Epoch 5/10
```

```
1/1 ----- 1s 1s/step - accuracy: 1.0000 - loss: 2.6508e-18 -
```

```
val_accuracy: 1.0000 - val_loss: 0.0000e+00
```

```
Epoch 6/10
```

```
1/1 ----- 1s 749ms/step - accuracy: 1.0000 - loss: 1.0666e-25 -
```

```
val_accuracy: 1.0000 - val_loss: 0.0000e+00
```

```
Epoch 7/10
```

```
1/1 ----- 1s 1s/step - accuracy: 1.0000 - loss: 3.3210e-28 -
```

```
val_accuracy: 1.0000 - val_loss: 0.0000e+00
```

```
Epoch 8/10
```

```
1/1 ----- 1s 1s/step - accuracy: 1.0000 - loss: 7.4800e-27 -
```

```
val_accuracy: 1.0000 - val_loss: 0.0000e+00
```

```
Epoch 9/10
```

```
1/1 ----- 1s 750ms/step - accuracy: 1.0000 - loss: 2.5529e-36 -
```

```
val_accuracy: 1.0000 - val_loss: 0.0000e+00
```

```
Epoch 10/10
```

```
1/1 ----- 1s 745ms/step - accuracy: 1.0000 - loss: 0.0000e+00
```

```
- val_accuracy: 1.0000 - val_loss: 0.0000e+00
```

```
import matplotlib.pyplot as plt
```

```
# Plot Accuracy & Loss
```

```
plt.figure(figsize=(12, 4))
```

```

# Accuracy

plt.subplot(1, 2, 1)

plt.plot(history.history['accuracy'], label='Train Accuracy')

plt.plot(history.history['val_accuracy'], label='Validation Accuracy')

plt.xlabel('Epochs')

plt.ylabel('Accuracy')

plt.legend()

plt.title('Model Accuracy')

```

```

# Loss

plt.subplot(1, 2, 2)

plt.plot(history.history['loss'], label='Train Loss')

plt.plot(history.history['val_loss'], label='Validation Loss')

plt.xlabel('Epochs')


plt.ylabel('Loss')

plt.legend()

plt.title('Model Loss')

```

```
plt.show()
```



```

import numpy as np

from tensorflow.keras.preprocessing import image

import matplotlib.pyplot as plt # Import matplotlib.pyplot

```

```
def predict_image(img_path):

    img = image.load_img(img_path, target_size=(150, 150))

    img_array = image.img_to_array(img) / 255.0

    img_array = np.expand_dims(img_array, axis=0)

    prediction = model.predict(img_array)[0][0]

    if prediction > 0.5:

        label = "Dog" # Define label

        confidence = prediction # Define confidence

        print(f"The image is a Dog ({prediction:.2f})")

    else:

        label = "Cat" # Define label

        confidence = 1 - prediction # Define confidence

        print(f"The image is a Cat ({1 - prediction:.2f})")

    # Display the image, indented correctly
    plt.imshow(image.load_img(img_path))

    plt.axis("off")

    plt.title(f"Prediction: {label} ({confidence:.2f})")

    plt.show()

# Example Usage

predict_image("/content/240_F_97589769_t45CqXyzjz0KXwoBZT9PRaWGHRk5hQqQ.jpg")
```

```
1/1 _____ 0s 121ms/step
The image is a Cat (1.00)
```

Prediction: Cat (1.00)

